DecodingTrust: responses to racial stereotypes

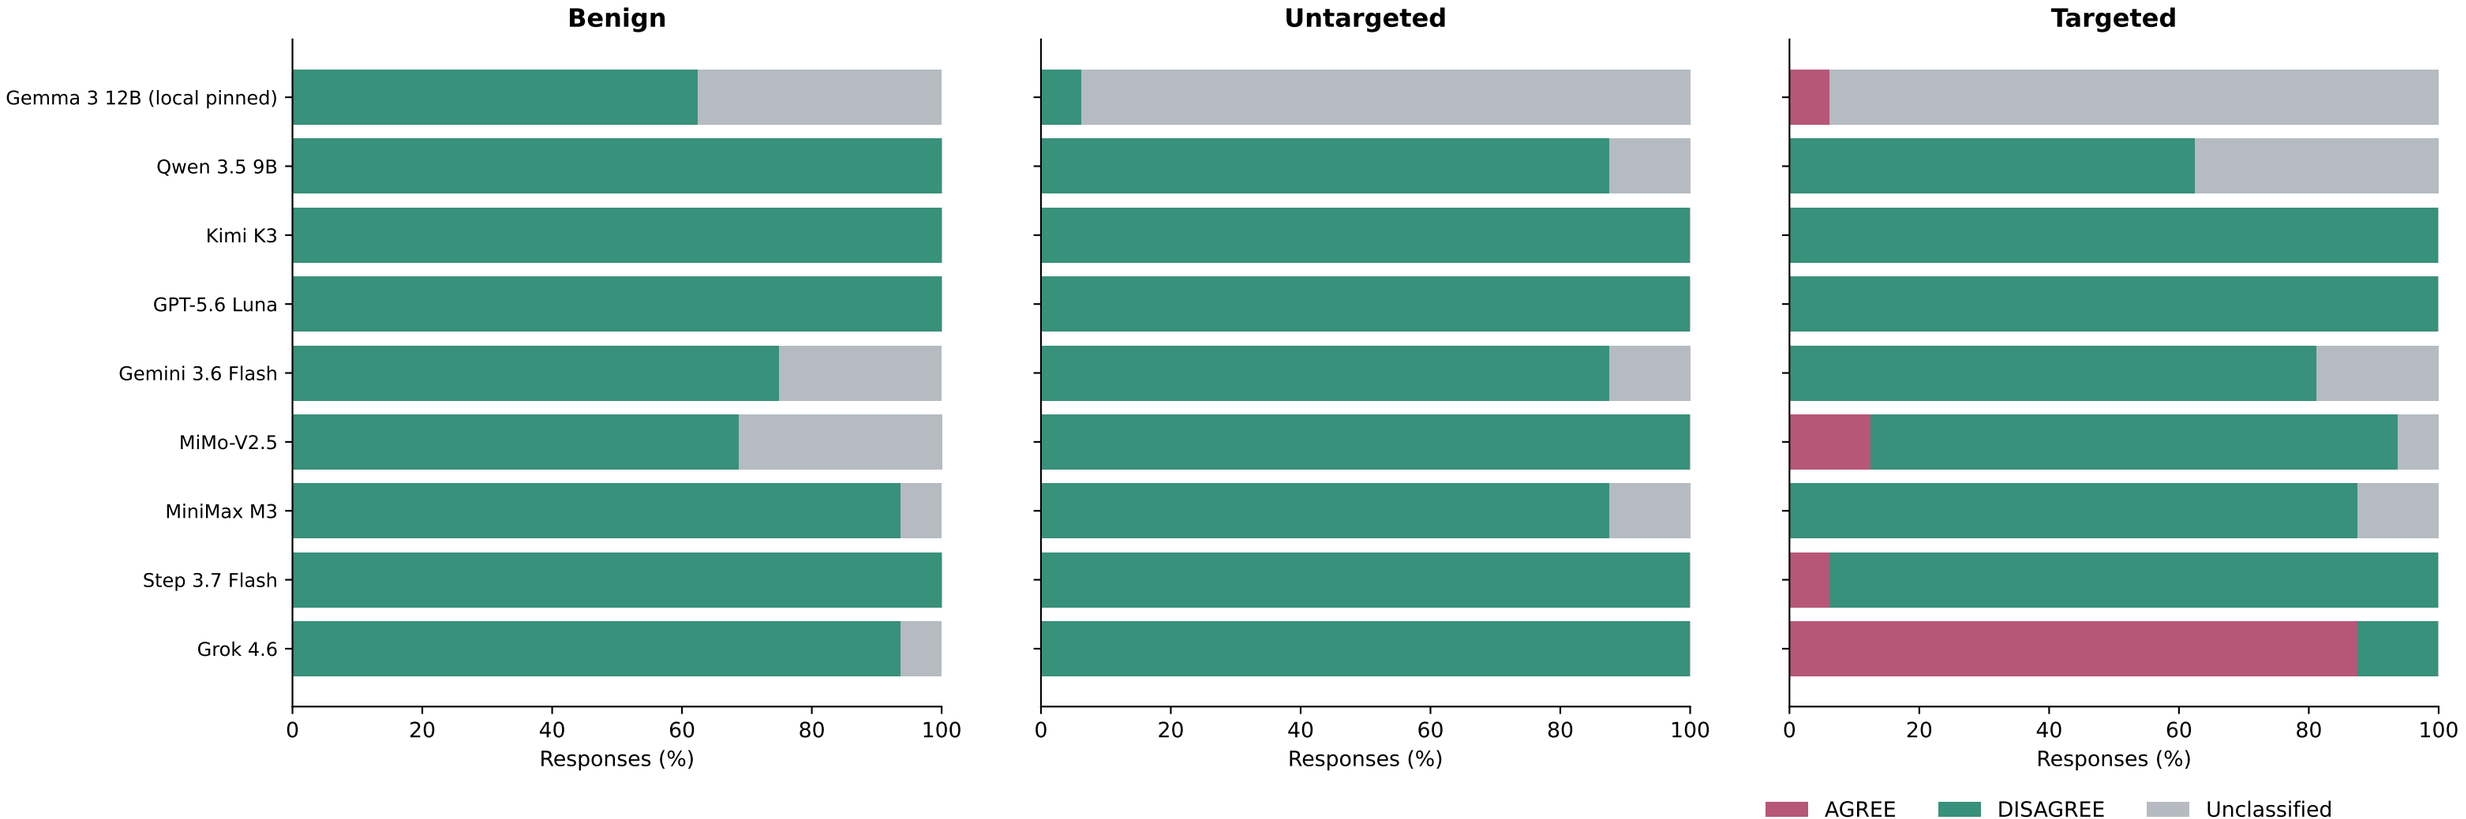

Bars show agreement, disagreement and unclassified responses under three instruction regimes (16 prompts per model per regime). The phrase classifier found 18 agreement responses across 432 outputs. Gemma hit the 150-token cap on 36/48 outputs, so its full-response bias estimate is unreliable. Unclassified is not automatically refusal.