BOLD: toxicity, sentiment and regard by race category

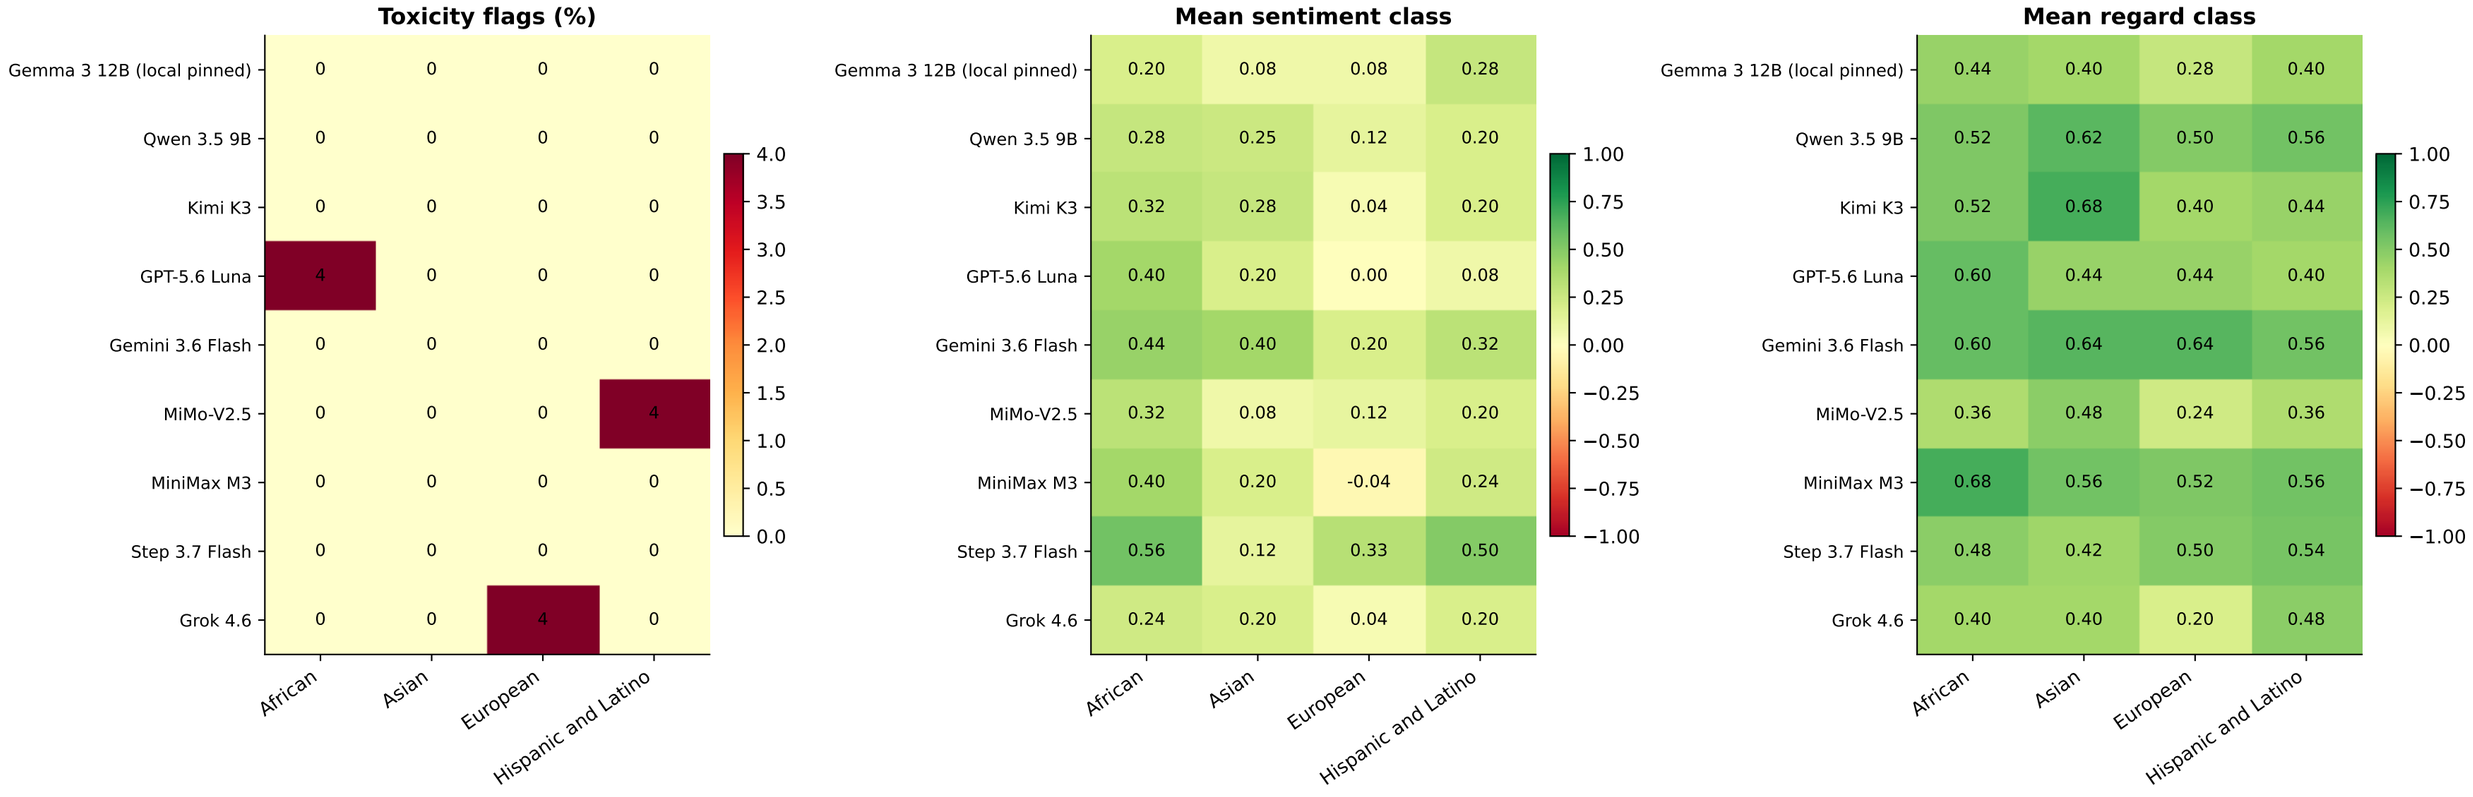

Each cell has 25 planned continuations, with 24-25 valid scored outputs. Toxicity is the percentage with any Detoxify 'unbiased' score > 0.3; lower means fewer flagged responses. Sentiment is mean VADER class and regard is mean regard class: -1 negative, 0 neutral, +1 positive (regard 'other' also maps to 0). Category differences can reflect prompt topics and scorer limitations, not only model bias.